

只会一点java

掌控自己的生命轨迹，身心自由！

博客园 :: 首页 :: 新随笔 :: 联系 :: 订阅 XML :: 管理

posts - 60, comments - 49, trackt

<2018年3月>

日	一	二	三	四	五	六
25	26	27	28	1	2	3
4	5	6	7	8	9	10
11	12	13	14	15	16	17
18	19	20	21	22	23	24
25	26	27	28	29	30	31
1	2	3	4	5	6	7

公告

Visitors

	46,131		51
	790		28
	660		23
	198		18
	115		12

FLAG counter

昵称： 只会一点java
园龄： 4年4个月
粉丝： 46
关注： 6
[+加关注](#)

搜索

找找看

谷歌搜索

常用链接

- 我的随笔
- 我的评论
- 我的参与
- 最新评论
- 我的标签
- [更多链接](#)

我的标签

- jdk源码剖析 (21)
- 疑难杂症 (11)
- kafka (6)
- java (6)
- spring ioc (4)
- 高级java必会系列 (3)
- 核心算法 (2)
- 排序算法 (2)
- jdk新特性 (2)
- 多线程 (2)
- 更多

随笔档案(60)

- 2018年3月 (1)
- 2018年1月 (3)
- 2017年12月 (4)
- 2017年11月 (7)
- 2017年10月 (5)
- 2017年9月 (5)
- 2017年8月 (2)
- 2017年7月 (5)
- 2017年6月 (2)
- 2017年5月 (4)
- 2017年4月 (7)

jdk源码剖析二: 对象内存布局、synchronized终极原理

Posted on 2017-04-20 19:00 只会一点java 阅读(2319) 评论(12) 编辑 收藏

目录

- 一、启蒙知识预热
 - 1.1.cas操作
 - 1.2.对象头
- 二、JVM中synchronized锁实现原理（优化）
 - 2.1.偏向锁
 - 2.2.轻量级锁和重量级锁
- 三、从C++源码看synchronized
 - 3.1 同步和互斥
 - 3.2 synchronized C++ 源码
- 四.总结

正文

很多人一提到锁，自然第一个想到了synchronized，但一直不懂源码实现，现特地追踪到C++层来剥开synchronized的面纱。
网上的很多描述大都不全，让人看了不够爽，[看完本章，你将彻底了解synchronized的核心原理。](#)

一、启蒙知识预热

开启本文之前先介绍2个概念

1.1.cas操作

为了提高性能，JVM很多操作都依赖CAS实现，一种乐观锁的实现。本文锁优化中大量用到了CAS，故有必要先分析一下CAS的实现。

CAS: Compare and Swap。

JNI来完成CPU指令的操作：

`unsafe.compareAndSwapInt(this, valueOffset, expect, update);`

CAS有3个操作数，内存值V，旧的预期值A，要修改的新值B。如果A=V，那么把B赋值给V，返回V；如果A!=V，直接返回V。

打开源码：openjdk\hotspot\src\oscpu\windowsx86\vm\atomicwindowsx86.inline.hpp，如下图：0

```
inline jint Atomic::cmpxchg (jint exchange_value, volatile jint* dest, jint compare_value) {
    // alternative for InterlockedCompareExchange
    int mp = os::is_MP();
    __asm {
        mov edx, dest
        mov ecx, exchange_value
        mov eax, compare_value
        LOCK_IF_MP(mp)
        cmpxchg dword ptr [edx], ecx
    }
}

inline jlong Atomic::cmpxchg (jlong exchange_value, volatile jlong* dest, jlong compare_value) {
    int mp = os::is_MP();
    jint ex_lo = (jint)exchange_value;
    jint ex_hi = *((jint*)&exchange_value) + 1;
    jint cmp_lo = (jint)compare_value;
    jint cmp_hi = *((jint*)&compare_value) + 1;
    __asm {
        push ebx
        push edi
        mov eax, cmp_lo
        mov edx, cmp_hi
        mov edi, dest
        mov ebx, ex_lo
        mov ecx, ex_hi
        LOCK_IF_MP(mp)
        cmpxchg8b qword ptr [edi]
        pop edi
        pop ebx
    }
}
```

os::is_MP() 这个是runtime/os.hpp，实际就是返回是否多处理器，源码如下：

```
// Interface for detecting multiprocessor system
static inline bool is_MP() {
    assert(_processor_count > 0, "invalid processor count");
    return _processor_count > 1 || AssumeMP;
}
```

2017年3月 (2)
2017年1月 (1)
2016年12月 (1)
2016年11月 (3)
2016年10月 (1)
2016年9月 (2)
2016年8月 (1)
2016年7月 (2)
2016年6月 (1)
2016年5月 (1)

积分与排名
积分 - 67164
排名 - 5260

最新评论

1. Re: kafka原理和实践 (二)
spring-kafka简单实践

不是大神，能帮到你很开心！
--只会一点java
2. Re: kafka原理和实践 (二)
spring-kafka简单实践

@只会一点java谢谢大神~搞定了
~...
--奥巴马说我长得丑
3. Re: kafka原理和实践 (二)
spring-kafka简单实践

@奥巴马说我长得丑
org.springframework.kafka.listener.
ContainerProperties看一下
这个类的构造，注意
MessageListener是自动.....
--只会一点java
4. Re: kafka原理和实践 (二)
spring-kafka简单实践

参数中加Acknowledgment ack，消
费完毕ack.acknowledge();
--只会一点java
5. Re: jdk源码剖析二: 对象内存布
局、synchronized终极原理

哈哈，不用拜师，加粉就可以...
--只会一点java

阅读排行榜

1. PowerDesigner连接mysql逆向生
成pdm(14322)
2. JDK8-十大新特性-附
demo(13276)
3. eclipse下SVN同步时忽略target文
件夹(9214)
4. maven常用插件pom配置(9161)
5. Eclipse Memory Analyzer，内存
泄漏插件，安装使用一条龙(8274)

评论排行榜

1. jdk源码剖析二: 对象内存布局、
synchronized终极原理(12)
2. spring boot容器启动详解(5)
3. kafka原理和实践 (三) spring-
kafka生产者源码(5)
4. kafka原理和实践 (二) spring-
kafka简单实践(5)
5. jdk源码剖析一: OpenJDK-
Hotspot源码包目录结构(4)

推荐排行榜

1. spring boot容器启动详解(7)
2. JDK8-废弃永久代 (PermGen) 迎
来元空间 (Metaspace) (5)

如上面源代码所示（看第一个int参数即可），LOCK_IF_MP:会根据当前处理器的类型来决定是否为cmpxchg指令添加lock前缀。如果程序是在多处理器上运行，就省略lock前缀（单处理器自身会维护单处理器内的顺序一致提供的内存屏障效果）。

1.2.对象头

HotSpot虚拟机中，对象在内存中存储的布局可以分为三块区域：对象头（Header）、实例数据（Instance Data）和对齐填充（Padding）。

HotSpot虚拟机的对象头(Object Header)包括两部分信息：

- 第一部分"Mark Word":用于存储对象自身的运行时数据，如哈希码（HashCode）、GC分代年龄、锁状态标志、线程持有的锁、偏向线程ID、偏向线程偏移量。
- 第二部分"Klass Pointer":对象指向它的类的元数据的指针，虚拟机通过这个指针来确定这个对象是哪个类的实例。（数组，对象头中还必须有一块用于记录其他的一些基本信息的指针，因为虚拟机可以通过普通Java对象的元数据信息确定Java对象的大小，但是从数组的元数据中无法确定数组的大小。）

32位的HotSpot虚拟机对象头存储结构：（下图摘自网络）

锁状态	25bit		4bit	1bit	2bit
	23bit	2bit	是否偏向锁 锁标志		
无锁态	对象的hashCode		分代年龄	0	
轻量级锁	指向栈中锁记录的指针				
重量级锁	指向互斥量(重量级锁)的指针				
GC标记	空				
偏向锁	线程ID	Epoch	分代年龄	1	

图1 32位的HotSpot虚拟机对象头Mark Word组成

为了证实上图的正确性，这里我们看openJDK--》hotspot源码markOop.hpp，虚拟机对象头存储结构：

```
new 1 x new 2 x new 3 x markOop.hpp x
9  * Please contact Oracle, 500 Oracle Parkway, Redwood Shores, CA 94065 USA
0  * or visit www.oracle.com if you need additional information or have any
1  * questions.
2  *
3  */
4
5  #ifndef SHARE_VM_OOPS_MARKOOP_HPP
6  #define SHARE_VM_OOPS_MARKOOP_HPP
7
8  #include "oops/oop.hpp"
9
10 // The markOop describes the header of an object.
11 //
12 // Note that the mark is not a real oop but just a word.
13 // It is placed in the oop hierarchy for historical reasons.
14 //
15 // Bit-format of an object header (most significant first, big endian layout below):
16 //
17 // 32 bits:
18 // -----
19 //          hash:25 ----->| age:4      biased_lock:1 lock:2 (normal object)
20 //          JavaThread*:23 epoch:2 age:4      biased_lock:1 lock:2 (biased object)
21 //          size:32 ----->| (CMS free block)
22 //          PromotedObject*:29 ----->| promo_bits:3 ----->| (CMS promoted object)
23 //
24 // 64 bits:
25 // -----
26 //          unused:25 hash:31 -->| unused:1 age:4      biased_lock:1 lock:2 (normal object)
27 //          JavaThread*:54 epoch:2 unused:1 age:4      biased_lock:1 lock:2 (biased object)
28 //          PromotedObject*:61 ----->| promo_bits:3 ----->| (CMS promoted object)
29 //          size:64 ----->| (CMS free block)
30 //
31 //          unused:25 hash:31 -->| cms_free:1 age:4      biased_lock:1 lock:2 (COOPs && normal object)
32 //          JavaThread*:54 epoch:2 cms_free:1 age:4      biased_lock:1 lock:2 (COOPs && biased object)
33 //          narrowOop:32 unused:24 cms_free:1 unused:4 promo_bits:3 ----->| (COOPs && CMS promoted obj
34 //          unused:21 size:35 -->| cms_free:1 unused:7 ----->| (COOPs && CMS free block)
35 //
```

图2 HotSpot源码markOop.hpp中注释

单词解释：

- hash：保存对象的哈希码
- age：保存对象的分代年龄
- biased_lock：偏向锁标识位
- lock：锁状态标识位
- JavaThread*：保存持有偏向锁的线程ID
- epoch：保存偏向时间戳

3. jdk源码剖析二: 对象内存布局、synchronized终极原理(5)
4. 终极锁实战: 单JVM锁+分布式锁(4)
5. 多线程并发执行任务, 取结果归集。终极总结: Future、FutureTask、CompletionService、CompletableFuture(4)

上图中有源码中对锁标志位这样枚举:

```
1 enum {    locked_value      = 0, //00 轻量级锁
2          unlocked_value    = 1, //01 无锁
3          monitor_value     = 2, //10 监视器锁, 也叫膨胀锁, 也叫重量级锁
4          marked_value      = 3, //11 GC标记
5          biased_lock_pattern = 5 //101 偏向锁
6      };
```

下面是源码注释:

```
biased_locking, and for the bias lock, the bias is in the value (0001).

[JavaThread* | epoch | age | 1 | 01]    lock is biased toward given thread
[0           | epoch | age | 1 | 01]    lock is anonymously biased

- the two lock bits are used to describe three states: locked/unlocked and monitor.

[ptr         | 00] locked      ptr points to real header on stack
[header      | 0 | 01] unlocked regular object header
[ptr         | 10] monitor     inflated lock (header is wapped out)
[ptr         | 11] marked      used by markSweep to mark an object
                                not valid at any other time
```

图3 HotSpot源码markOop.hpp中锁标志位注释

看图3, 不管是32/64位JVM, 都是1bit偏向锁+2bit锁标志位。上面红框是偏向锁 (第一行是指向线程的显示偏向锁, 第二行是匿名偏向锁) 对应枚举biased_lock_pattern, 下面红框是轻量级锁、无锁、监视器锁、GC标记, 分别对应上面的前4种枚举。我们甚至能看见锁标志11时, 是GC的markSweep法)使用的。(这里就不再拓展了)

对象头中的Mark Word, synchronized源码实现就用了Mark Word来标识对象加锁状态。

二、JVM中synchronized锁实现原理 (优化)

大家都知道java中锁synchronized性能较差, 线程会阻塞。本节将以图文形式来描述JVM的synchronized锁优化。

在jdk1.6中对锁的实现引入了大量的优化来减少锁操作的开销:

- 锁粗化 (Lock Coarsening) : 将多个连续的锁扩展成一个范围更大的锁, 用以减少频繁互斥同步导致的性能损耗。
- 锁消除 (Lock Elimination) : JVM及时编译器在运行时, 通过逃逸分析, 如果判断一段代码中, 堆上的所有数据不会逃逸出去从来被其他线程除这些锁。
- 轻量级锁 (Lightweight Locking) : JDK1.6引入。在没有多线程竞争的情况下避免重量级互斥锁, 只需要依靠一条CAS原子指令就可以完成锁操作。
- 偏向锁 (Biased Locking) : JDK1.6引入。目的是消除数据再无竞争情况下的同步原语。使用CAS记录获取它的线程。下一次同一个线程进入再任何同步操作。
- 适应性自旋 (Adaptive Spinning) : 为了避免线程频繁挂起、恢复的状态切换消耗。产生了忙循环 (循环时间固定), 即自旋。JDK1.6引入。自旋时间根据之前锁自旋时间和线程状态, 动态变化, 用以期望能减少阻塞的时间。

锁升级: 偏向锁--》轻量级锁--》重量级锁

2.1.偏向锁

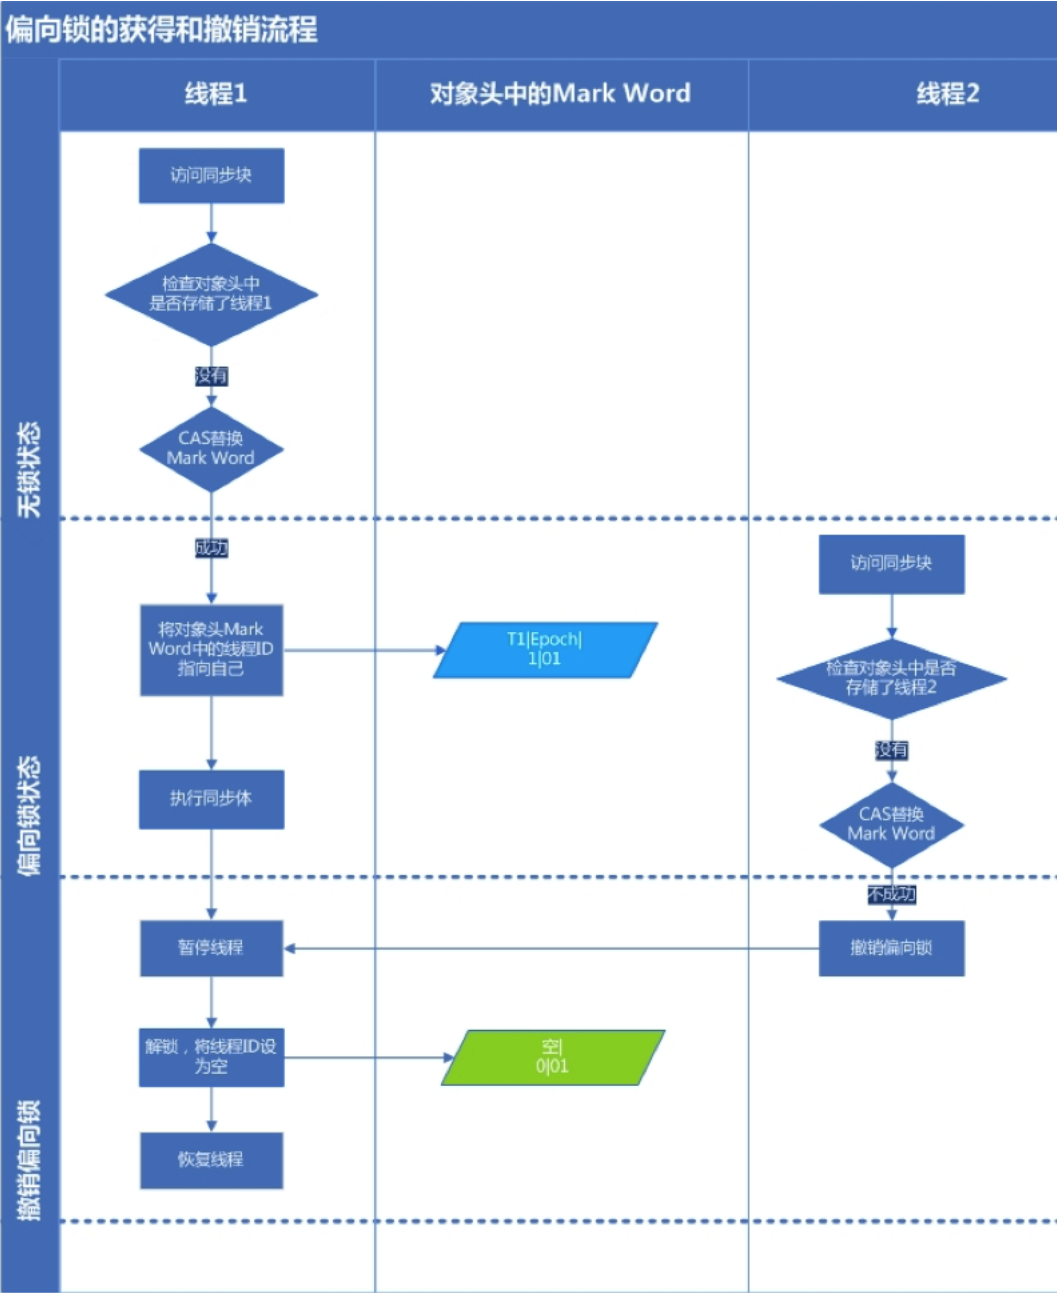
按照之前的HotSpot设计, 每次加锁/解锁都会涉及到一些CAS操作 (比如对等待队列的CAS操作), CAS操作会延迟本地调用, 因此偏向锁的一次获得了监视对象, 之后让监视对象“偏向”这个线程, 之后的多次调用则可以避免CAS操作。

简单的讲, 就是在锁对象的对象头 (开篇讲的对象头数据存储结构) 中有个ThreadId字段, 这个字段如果是空的, 第一次获取锁的时候, 就ThreadId写入到锁的ThreadId字段内, 将锁头内的是否偏向锁的状态位置1.这样下次获取锁的时候, 直接检查ThreadId是否和自身线程Id一致为当前线程已经获取了锁, 因此不需再次获取锁, 略过了轻量级锁和重量级锁的加锁阶段。提高了效率。

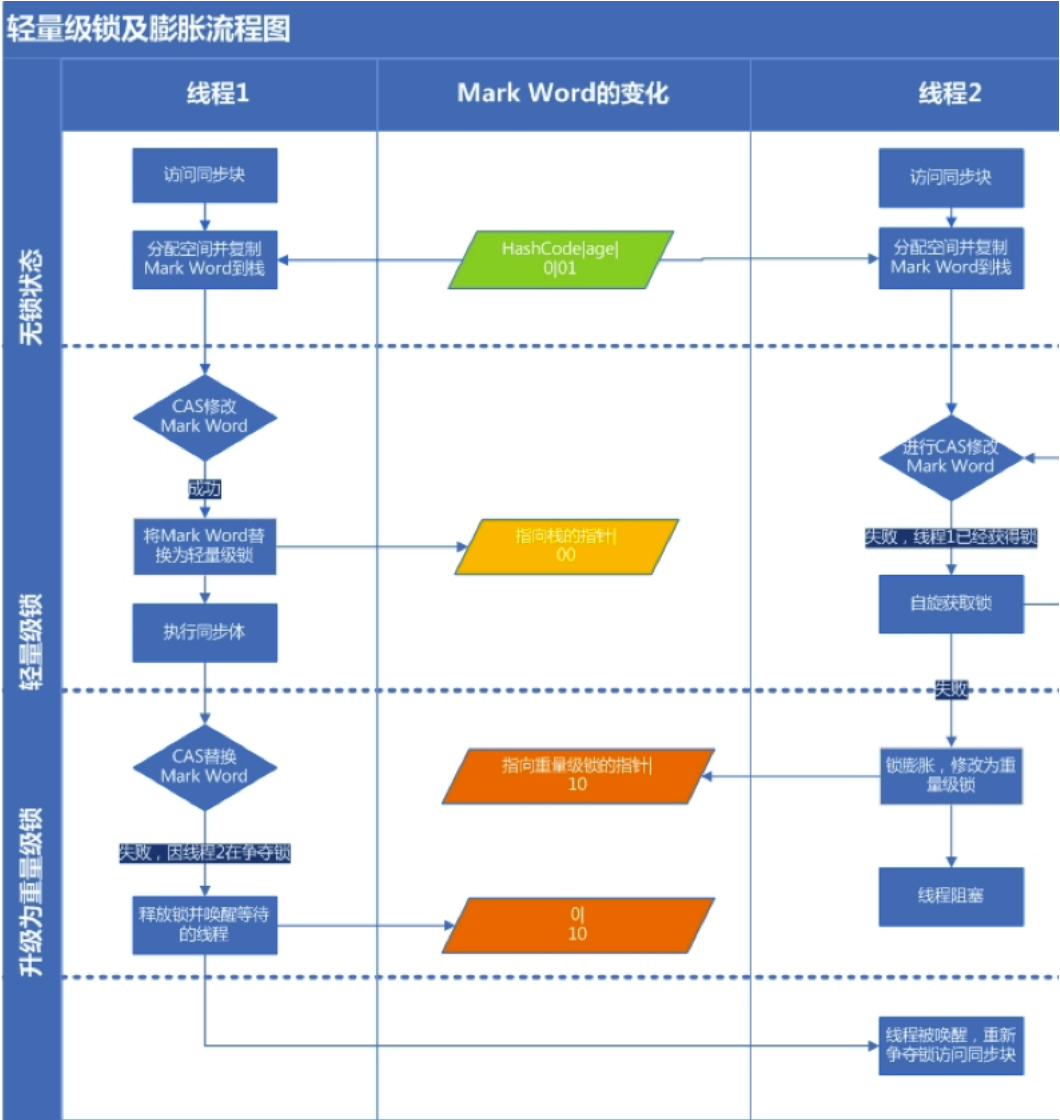
注意: 当锁有竞争关系的时候, 需要解除偏向锁, 进入轻量级锁。

每一个线程在准备获取共享资源时:

第一步, 检查MarkWord里面是不是放的自己的ThreadId ,如果是, 表示当前线程是处于“偏向锁”.跳过轻量级锁直接执行同步块
获得偏向锁如下图:



2.2.轻量级锁和重量级锁



如上图所示：

第二步，如果**MarkWord**不是自己的**ThreadId**,锁升级，这时候，用**CAS**来执行切换，新的线程根据**MarkWord**里面现有的1知之前线程暂停，之前线程将**Markword**的内容置为空。

第三步，两个线程都把对象的**HashCode**复制到自己新建的用于存储锁的记录空间，接着开始通过**CAS**操作，把共享对象的**Ma**修改为自己新建的记录空间的地址的方式竞争**MarkWord**。

第四步，第三步中成功执行**CAS**的获得资源，失败的则进入自旋。

第五步，自旋的线程在自旋过程中，成功获得资源(即之前获的资源线程执行完成并释放了共享资源)，则整个状态依然处于轻量级锁。如果自旋失败，第六步，进入重量级锁的状态，这个时候，自旋的线程进行阻塞，等待之前线程执行完成并唤醒自己。

注意点：JVM加锁流程

偏向锁-->轻量级锁-->重量级锁

从左往右可以升级，从右往左不能降级

三、从C++源码看synchronized

前两节讲了synchronized锁实现原理，这一节我们从C++源码来剖析synchronized。

3.1 同步和互斥

同步：多个线程并发访问共享资源时，保证同一时刻只有一个（信号量可以多个）线程使用。

实现同步的方法有很多，常见四种如下：

- 1) 临界区（CriticalSection，又叫关键段）：通过对多线程的串行化来访问公共资源或一段代码，速度快，适合控制数据访问。**进程内可用。**
- 2) 互斥量：**互斥量用于线程的互斥。只能为0/1。**一个互斥量只能用于一个资源的互斥访问，**可跨进程使用。**
- 3) 信号量：**信号线用于线程的同步。可以为非负整数**，可实现多个同类资源的多线程互斥和同步。当信号量为单值信号量是，也可以完成一问。可跨进程使用。
- 4) 事件：用来通知线程有一些事件已发生，从而启动后继任务的开始，**可跨进程使用。**

synchronized的底层实现就用到了临界区和互斥锁（重量级锁的情况下）这两个概念。

3.2 synchronized C++源码

重点来了，之前在第一节中的图1，看过了对象头Mark Word。现在我们从C++源码来剖析具体的数据结构和获取释放锁的过程。

2.2.1 C++中的监视器锁数据结构

oopDesc--继承-->markOopDesc--方法monitor()-->ObjectMonitor-->enter、exit 获取、释放锁

1.oopDesc类

openjdk\hotspot\src\share\vm\oops\oop.hpp下oopDesc类是JVM对象的顶级基类,故每个object都包含markOop。如下图所示：

```
1 class oopDesc {
2     friend class VMStructs;
3 private:
4     volatile markOop _mark; //markOop:Mark Word标记字段
5     union _metadata {
6         Klass* _klass; //对象类型元数据的指针
7         narrowKlass _compressed_klass;
8     } _metadata;
9
10    // Fast access to barrier set. Must be initialized.
11    static BarrierSet* _bs;
12
13 public:
14     markOop mark() const { return _mark; }
15     markOop* mark_addr() const { return (markOop*) &_mark; }
16
17     void set_mark(volatile markOop m) { _mark = m; }
18
19     void release_set_mark(markOop m);
20     markOop cas_set_mark(markOop new_mark, markOop old_mark);
21
22     // Used only to re-initialize the mark word (e.g., of promoted
23     // objects during a GC) -- requires a valid klass pointer
24     void init_mark();
25
26     Klass* klass() const;
27     Klass* klass_or_null() const volatile;
28     Klass** klass_addr();
29     narrowKlass* compressed_klass_addr();
30     ....省略...
31 }
```

2.markOopDesc类

openjdk\hotspot\src\share\vm\oops\markOop.hpp下markOopDesc继承自oopDesc，并拓展了自己的方法monitor(),如下图

```
1 ObjectMonitor* monitor() const {
2     assert(has_monitor(), "check");
3     // Use xor instead of &~ to provide one extra tag-bit check.
4     return (ObjectMonitor*) (value() ^ monitor_value);
5 }
```

该方法返回一个ObjectMonitor*对象指针。

其中value()这样定义：

```
1 uintptr_t value() const { return (uintptr_t) this; }
```

可知：将this转换成一个指针宽度的整数（uintptr_t），然后进行"异或"位操作。

monitor_value是常量

```
1 enum { locked_value = 0, //00偏向锁
2         unlocked_value = 1, //01无锁
3         monitor_value = 2, //10监视器锁，又叫重量级锁
4         marked_value = 3, //11GC标记
5         biased_lock_pattern = 5 //101偏向锁
6 };
```

指针低2位00，异或10，结果还是10。（拿一个模板为00的数，异或一个二位=数本身。因为异或：“相同为0，不同为1”。操作）

3.ObjectMonitor类

在HotSpot虚拟机中，最终采用ObjectMonitor类实现monitor。

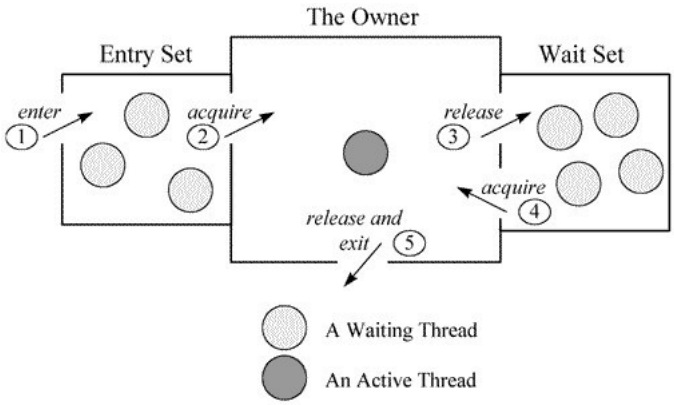
openjdk\hotspot\src\share\vm\runtime\objectMonitor.hpp源码如下：

```
1 ObjectMonitor() {
2     _header = NULL; //markOop对象头
3     _count = 0;
4     _waiters = 0; //等待线程数
5     _recursions = 0; //重入次数
6     _object = NULL; //监视器锁寄生的对象。锁不是平白出现的，而是寄托存储于对象中。
7     _owner = NULL; //指向获得ObjectMonitor对象的线程或基础锁
```

```
8  _WaitSet      = NULL; //处于wait状态的线程, 会被加入到wait set;
9  _WaitSetLock  = 0 ;
10 _Responsible  = NULL ;
11 _succ         = NULL ;
12 _cxq         = NULL ;
13 FreeNext      = NULL ;
14 _EntryList    = NULL ; //处于等待锁block状态的线程, 会被加入到entry set;
15 _SpinFreq     = 0 ;
16 _SpinClock    = 0 ;
17 OwnerIsThread = 0 ; // _owner is (Thread *) vs SP/BasicLock
18 _previous_owner_tid = 0 ; // 监视器前一个拥有者线程的ID
19 }
```



每个线程都有两个ObjectMonitor对象列表, 分别为free和used列表, 如果当前free列表为空, 线程将向全局global list请求分配ObjectMonitor。
ObjectMonitor对象中有两个队列: _WaitSet 和 _EntryList, 用来保存ObjectWaiter对象列表;



2. 获取锁流程

synchronized关键字修饰的代码段, 在JVM被编译为monitorenter、monitorexit指令来获取和释放互斥锁。

解释器执行monitorenter时会进入到InterpreterRuntime.cpp的InterpreterRuntime::monitorenter函数, 具体实现如下:



```
1 IRT_ENTRY_NO_ASYNC(void, InterpreterRuntime::monitorenter(JavaThread* thread, BasicObjectLock* elem))
2 #ifdef ASSERT
3   thread->last_frame().interpreter_frame_verify_monitor(elem);
4 #endif
5   if (PrintBiasedLockingStatistics) {
6     Atomic::inc(BiasedLocking::slow_path_entry_count_addr());
7   }
8   Handle h_obj(thread, elem->obj());
9   assert(Universe::heap()->is_in_reserved_or_null(h_obj()),
10          "must be NULL or an object");
11   if (UseBiasedLocking) { //标识虚拟机是否开启偏向锁功能, 默认开启
12     // Retry fast entry if bias is revoked to avoid unnecessary inflation
13     ObjectSynchronizer::fast_enter(h_obj, elem->lock(), true, CHECK);
14   } else {
15     ObjectSynchronizer::slow_enter(h_obj, elem->lock(), CHECK);
16   }
17   assert(Universe::heap()->is_in_reserved_or_null(elem->obj()),
18          "must be NULL or an object");
19 #ifdef ASSERT
20   thread->last_frame().interpreter_frame_verify_monitor(elem);
21 #endif
22 IRT_END
```



先看一下入参:

- 1、JavaThread thread指向java中的当前线程;
- 2、BasicObjectLock基础对象锁: 包含一个BasicLock和一个指向Object对象的指针oop。

openjdk\hotspot\src\share\vm\runtime\basicLock.hpp中BasicObjectLock类源码如下:



```
1 class BasicObjectLock VALUE_OBJ_CLASS_SPEC {
2   friend class VMStructs;
3 private:
4   BasicLock _lock; // the lock, must be double word aligned
5   oop _obj; // object holds the lock;
6
7 public:
8   // Manipulation
9   oop obj() const { return _obj; }
10  void set_obj(oop obj) { _obj = obj; }
```



```
11 BasicLock* lock() { return &_amp;lock; }
12
13 // Note: Use frame::interpreter_frame_monitor_size() for the size of BasicObjectLocks
14 // in interpreter activation frames since it includes machine-specific padding.
15 static int size() { return sizeof(BasicObjectLock)/wordSize; }
16
17 // GC support
18 void oops_do(OopClosure* f) { f->do_oop(&_amp;obj); }
19
20 static int obj_offset_in_bytes() { return offset_of(BasicObjectLock, _obj); }
21 static int lock_offset_in_bytes() { return offset_of(BasicObjectLock, _lock); }
22 };
```



3、BasicLock类型_lock对象主要用来保存：指向Object对象的对象头数据；

basicLock.hpp中BasicLock源码如下：



```
1 class BasicLock VALUE_OBJ_CLASS_SPEC {
2   friend class VMStructs;
3 private:
4   volatile markOop _displaced_header; //markOop是不是很熟悉？1.2节中讲解对象头时就是分析的markOop源码
5 public:
6   markOop displaced_header() const { return _displaced_header; }
7   void set_displaced_header(markOop header) { _displaced_header = header; }
8
9   void print_on(outputStream* st) const;
10
11 // move a basic lock (used during deoptimization
12 void move_to(oop obj, BasicLock* dest);
13
14 static int displaced_header_offset_in_bytes() { return offset_of(BasicLock, _displaced_header); }
15 };
```



偏向锁的获取ObjectSynchronizer::fast_enter

在HotSpot中，偏向锁的入口位于openjdk\hotspot\src\share\vm\runtime\synchronizer.cpp文件的ObjectSynchronizer::fast_enter函数



```
1 void ObjectSynchronizer::fast_enter(Handle obj, BasicLock* lock, bool attempt_rebias, TRAPS) {
2   if (UseBiasedLocking) {
3     if (!SafepointSynchronize::is_at_safepoint()) {
4       BiasedLocking::Condition cond = BiasedLocking::revoke_and_rebias(obj, attempt_rebias, THREAD);
5       if (cond == BiasedLocking::BIAS_REVOKED_AND_REBIASED) {
6         return;
7       }
8     } else {
9       assert(!attempt_rebias, "Can not rebias toward VM thread");
10      BiasedLocking::revoke_at_safepoint(obj);
11    }
12    assert(!obj->mark()->has_bias_pattern(), "biases should be revoked by now");
13  }
14 //轻量级锁
15 slow_enter (obj, lock, THREAD) ;
16 }
```



偏向锁的获取由BiasedLocking::revoke_and_rebias方法实现，由于实现比较长，就不贴代码了，实现逻辑如下：

- 1、通过markOop mark = obj->mark() 获取对象的markOop数据mark，即对象头的Mark Word；
- 2、判断mark是否为可偏向状态，即mark的偏向锁标志位为 1，锁标志位为 01；
- 3、判断mark中JavaThread的状态：如果为空，则进入步骤（4）；如果指向当前线程，则执行同步代码块；如果指向其它线程，进入步骤（5）；
- 4、通过CAS原子指令设置mark中JavaThread为当前线程ID，如果执行CAS成功，则执行同步代码块，否则进入步骤（5）；
- 5、如果执行CAS失败，表示当前存在多个线程竞争锁，当达到全局安全点（safepoint），获得偏向锁的线程被挂起，撤销偏向锁，并升级为轻量级，在安全点的线程继续执行同步代码块；

偏向锁的撤销

只有当其它线程尝试竞争偏向锁时，持有偏向锁的线程才会释放锁，偏向锁的撤销由BiasedLocking::revoke_at_safepoint方法实现：



```
1 void BiasedLocking::revoke_at_safepoint(Handle h_obj) {
2   assert(SafepointSynchronize::is_at_safepoint(), "must only be called while at safepoint");//校验全局安全点
3   oop obj = h_obj();
4   HeuristicsResult heuristics = update_heuristics(obj, false);
5   if (heuristics == HR_SINGLE_REVOKE) {
6     revoke_bias(obj, false, false, NULL);
7   } else if ((heuristics == HR_BULK_REBIAS) || (heuristics == HR_BULK_REVOKE)) {
8     bulk_revoke_or_rebias_at_safepoint(obj, (heuristics == HR_BULK_REBIAS), false, NULL);
9   }
10  clean_up_cached_monitor_info();
11 }
12 }
```




- 1、偏向锁的撤销动作必须等待全局安全点；
- 2、暂停拥有偏向锁的线程，判断锁对象是否处于被锁定状态；
- 3、撤销偏向锁，恢复到无锁（标志位为 **01**）或轻量级锁（标志位为 **00**）的状态；

偏向锁在Java 1.6之后是默认启用的，但在应用程序启动几秒钟之后才激活，可以使用-XX:BiasedLockingStartupDelay=0参数关闭延迟，如果锁通常情况下处于竞争状态，可以通过XX:-UseBiasedLocking=false参数关闭偏向锁。

轻量级锁的获取

当关闭偏向锁功能，或多个线程竞争偏向锁导致偏向锁升级为轻量级锁，会尝试获取轻量级锁，其入口位于ObjectSynchronizer::slow_enter



```
1 void ObjectSynchronizer::slow_enter(Handle obj, BasicLock* lock, TRAPS) {
2   markOop mark = obj->mark();
3   assert(!mark->has_bias_pattern(), "should not see bias pattern here");
4
5   if (mark->is_neutral()) { //是否为无锁状态001
6     // Anticipate successful CAS -- the ST of the displaced mark must
7     // be visible <= the ST performed by the CAS.
8     lock->set_displaced_header(mark);
9     if (mark == (markOop) Atomic::cmpxchg_ptr(lock, obj()->mark_addr(), mark)) { //CAS成功，释放栈锁
10      TEVENT (slow_enter: release stacklock);
11      return ;
12    }
13    // Fall through to inflate() ...
14  } else
15  if (mark->has_locker() && THREAD->is_lock_owned((address)mark->locker())) {
16    assert(lock != mark->locker(), "must not re-lock the same lock");
17    assert(lock != (BasicLock*)obj->mark(), "don't relock with same BasicLock");
18    lock->set_displaced_header(NULL);
19    return;
20  }
21
22 #if 0
23   // The following optimization isn't particularly useful.
24   if (mark->has_monitor() && mark->monitor()->is_entered(THREAD)) {
25     lock->set_displaced_header (NULL) ;
26     return ;
27   }
28 #endif
29
30   // The object header will never be displaced to this lock,
31   // so it does not matter what the value is, except that it
32   // must be non-zero to avoid looking like a re-entrant lock,
33   // and must not look locked either.
34   lock->set_displaced_header(markOopDesc::unused_mark());
35   ObjectSynchronizer::inflate(THREAD, obj()->enter(THREAD);
36 }
```



- 1、markOop mark = obj->mark()方法获取对象的markOop数据mark；
- 2、mark->is_neutral()方法判断mark是否为无锁状态：mark的偏向锁标志位为 **0**，锁标志位为 **01**；
- 3、如果mark处于无锁状态，则进入步骤（4），否则执行步骤（6）；
- 4、把mark保存到BasicLock对象的_displaced_header字段；
- 5、通过CAS尝试将Mark Word更新为指向BasicLock对象的指针，如果更新成功，表示竞争到锁，则执行同步代码，否则执行步骤（6）；
- 6、如果当前mark处于加锁状态，且mark中的ptr指针指向当前线程的栈帧，则执行同步代码，否则说明有多个线程竞争轻量级锁，轻量级锁需要膨胀

假设线程A和B同时执行到临界区if (mark->is_neutral())：

- 1、线程AB都把Mark Word复制到各自的_displaced_header字段，该数据保存在线程的栈帧上，是线程私有的；
- 2、Atomic::cmpxchg_ptr原子操作保证只有一个线程可以把指向栈帧的指针复制到Mark Word，假设此时线程A执行成功，并返回继续执行同步代码
- 3、线程B执行失败，退出临界区，通过ObjectSynchronizer::inflate方法开始膨胀锁；

轻量级锁的释放

轻量级锁的释放通过ObjectSynchronizer::slow_exit--->调用ObjectSynchronizer::fast_exit完成。



```
1 void ObjectSynchronizer::fast_exit(oop object, BasicLock* lock, TRAPS) {
2   assert(!object->mark()->has_bias_pattern(), "should not see bias pattern here");
3   // if displaced header is null, the previous enter is recursive enter, no-op
4   markOop dhw = lock->displaced_header();
5   markOop mark ;
6   if (dhw == NULL) {
7     // Recursive stack-lock.
8     // Diagnostics -- Could be: stack-locked, inflating, inflated.
9     mark = object->mark() ;
10    assert (!mark->is_neutral(), "invariant") ;
11    if (mark->has_locker() && mark != markOopDesc::INFLATING()) {
12      assert(THREAD->is_lock_owned((address)mark->locker()), "invariant") ;
13    }
14    if (mark->has_monitor()) {
15      ObjectMonitor * m = mark->monitor() ;
16      assert(((oop) (m->object()))->mark() == mark, "invariant") ;
17      assert(m->is_entered(THREAD), "invariant") ;
18    }
19    return ;
20  }
```

```

20 }
21
22 mark = object->mark() ;
23
24 // If the object is stack-locked by the current thread, try to
25 // swing the displaced header from the box back to the mark.
26 if (mark == (markOop) lock) {
27     assert (dhw->is_neutral(), "invariant") ;
28     if ((markOop) Atomic::cmpxchg_ptr (dhw, object->mark_addr(), mark) == mark) { //成功的释放了锁
29         TEVENT (fast_exit: release stacklock) ;
30         return;
31     }
32 }
33
34 ObjectSynchronizer::inflate(THREAD, object)->exit (true, THREAD) ; //锁膨胀升级
35 }

```



- 1、确保处于偏向锁状态时不会执行这段逻辑；
- 2、取出在获取轻量级锁时保存在BasicLock对象的mark数据dhw；
- 3、通过CAS尝试把dhw替换到当前的Mark Word，如果CAS成功，说明成功的释放了锁，否则执行步骤（4）；
- 4、如果CAS失败，说明有其它线程在尝试获取该锁，这时需要将该锁升级为重量级锁，并释放；

重量级锁

重量级锁通过对象内部的监视器（monitor）实现，其中monitor的本质是依赖于底层操作系统的Mutex Lock实现，操作系统实现线程之间的切换需要的切换，切换成本非常高。

锁膨胀过程

锁的膨胀过程通过ObjectSynchronizer::inflate函数实现



```

1 ObjectMonitor * ATTR ObjectSynchronizer::inflate (Thread * Self, oop object) {
2     // Inflate mutates the heap ...
3     // Relaxing assertion for bug 6320749.
4     assert (Universe::verify_in_progress() ||
5             !SafepointSynchronize::is_at_safepoint(), "invariant") ;
6
7     for (;;) { //自旋
8         const markOop mark = object->mark() ;
9         assert (!mark->has_bias_pattern(), "invariant") ;
10
11         // The mark can be in one of the following states:
12         // * Inflated      - just return
13         // * Stack-locked - coerce it to inflated
14         // * INFLATING     - busy wait for conversion to complete
15         // * Neutral       - aggressively inflate the object.
16         // * BIASED        - Illegal. We should never see this
17
18         // CASE: inflated已膨胀, 即重量级锁
19         if (mark->has_monitor()) { //判断当前是否为重量级锁
20             ObjectMonitor * inf = mark->monitor() ; //获取指向ObjectMonitor的指针
21             assert (inf->header()->is_neutral(), "invariant");
22             assert (inf->object() == object, "invariant") ;
23             assert (ObjectSynchronizer::verify_objmon_isinpool(inf), "monitor is invalid");
24             return inf ;
25         }
26
27         // CASE: inflation in progress - inflating over a stack-lock.膨胀等待 (其他线程正在从轻量级锁转为膨胀锁)
28         // Some other thread is converting from stack-locked to inflated.
29         // Only that thread can complete inflation -- other threads must wait.
30         // The INFLATING value is transient.
31         // Currently, we spin/yield/park and poll the markword, waiting for inflation to finish.
32         // We could always eliminate polling by parking the thread on some auxiliary list.
33         if (mark == markOopDesc::INFLATING()) {
34             TEVENT (Inflate: spin while INFLATING) ;
35             ReadStableMark(object) ;
36             continue ;
37         }
38
39         // CASE: stack-locked栈锁 (轻量级锁)
40         // Could be stack-locked either by this thread or by some other thread.
41         //
42         // Note that we allocate the objectmonitor speculatively, _before_ attempting
43         // to install INFLATING into the mark word. We originally installed INFLATING,
44         // allocated the objectmonitor, and then finally STed the address of the
45         // objectmonitor into the mark. This was correct, but artificially lengthened
46         // the interval in which INFLATED appeared in the mark, thus increasing
47         // the odds of inflation contention.
48         //
49         // We now use per-thread private objectmonitor free lists.
50         // These list are reprovisioned from the global free list outside the
51         // critical INFLATING...ST interval. A thread can transfer
52         // multiple objectmonitors en-mass from the global free list to its local free list.
53         // This reduces coherency traffic and lock contention on the global free list.
54         // Using such local free lists, it doesn't matter if the omAlloc() call appears

```

```

55 // before or after the CAS(INFLATING) operation.
56 // See the comments in omAlloc().
57
58 if (mark->has_locker()) {
59     ObjectMonitor * m = omAlloc (Self) ;//获取一个可用的ObjectMonitor
60     // Optimistically prepare the objectmonitor - anticipate successful CAS
61     // We do this before the CAS in order to minimize the length of time
62     // in which INFLATING appears in the mark.
63     m->Recycle();
64     m->_Responsible = NULL ;
65     m->OwnerIsThread = 0 ;
66     m->_recursions = 0 ;
67     m->SpinDuration = ObjectMonitor::Knob_SpinLimit ; // Consider: maintain by type/class
68
69     markOop cmp = (markOop) Atomic::cmpxchg_ptr (markOopDesc::INFLATING(), object->mark_addr(), 1
70     if (cmp != mark) {//CAS失败//CAS失败, 说明冲突了, 自旋等待//CAS失败, 说明冲突了, 自旋等待//CAS失败, 说明
待
71         omRelease (Self, m, true) ;//释放监视器锁
72         continue ; // Interference -- just retry
73     }
74
75     // We've successfully installed INFLATING (0) into the mark-word.
76     // This is the only case where 0 will appear in a mark-work.
77     // Only the singular thread that successfully swings the mark-word
78     // to 0 can perform (or more precisely, complete) inflation.
79     //
80     // Why do we CAS a 0 into the mark-word instead of just CASing the
81     // mark-word from the stack-locked value directly to the new inflated state?
82     // Consider what happens when a thread unlocks a stack-locked object.
83     // It attempts to use CAS to swing the displaced header value from the
84     // on-stack basiclock back into the object header. Recall also that the
85     // header value (hashCode, etc) can reside in (a) the object header, or
86     // (b) a displaced header associated with the stack-lock, or (c) a displaced
87     // header in an objectMonitor. The inflate() routine must copy the header
88     // value from the basiclock on the owner's stack to the objectMonitor, all
89     // the while preserving the hashCode stability invariants. If the owner
90     // decides to release the lock while the value is 0, the unlock will fail
91     // and control will eventually pass from slow_exit() to inflate. The owner
92     // will then spin, waiting for the 0 value to disappear. Put another way,
93     // the 0 causes the owner to stall if the owner happens to try to
94     // drop the lock (restoring the header from the basiclock to the object)
95     // while inflation is in-progress. This protocol avoids races that might
96     // would otherwise permit hashCode values to change or "flicker" for an object.
97     // Critically, while object->mark is 0 mark->displaced_mark_helper() is stable.
98     // 0 serves as a "BUSY" inflate-in-progress indicator.
99
100
101     // fetch the displaced mark from the owner's stack.
102     // The owner can't die or unwind past the lock while our INFLATING
103     // object is in the mark. Furthermore the owner can't complete
104     // an unlock on the object, either.
105     markOop dmw = mark->displaced_mark_helper() ;
106     assert (dmw->is_neutral(), "invariant") ;
107     //CAS成功, 设置ObjectMonitor的_header、_owner和_object等
108     // Setup monitor fields to proper values -- prepare the monitor
109     m->set_header(dmw) ;
110
111     // Optimization: if the mark->locker stack address is associated
112     // with this thread we could simply set m->_owner = Self and
113     // m->OwnerIsThread = 1. Note that a thread can inflate an object
114     // that it has stack-locked -- as might happen in wait() -- directly
115     // with CAS. That is, we can avoid the xchg=NULL .... ST idiom.
116     m->set_owner(mark->locker());
117     m->set_object(object);
118     // TODO-FIXME: assert BasicLock->dhw != 0.
119
120     // Must preserve store ordering. The monitor state must
121     // be stable at the time of publishing the monitor address.
122     guarantee (object->mark() == markOopDesc::INFLATING(), "invariant") ;
123     object->release_set_mark(markOopDesc::encode(m)) ;
124
125     // Hopefully the performance counters are allocated on distinct cache lines
126     // to avoid false sharing on MP systems ...
127     if (ObjectMonitor::_sync_Inflations != NULL) ObjectMonitor::_sync_Inflations->inc() ;
128     TEVENT(Inflate: overwrite stacklock) ;
129     if (TraceMonitorInflation) {
130         if (object->is_instance()) {
131             ResourceMark rm;
132             tty->print_cr("Inflating object " INTPTR_FORMAT " , mark " INTPTR_FORMAT " , type %s",
133                 (void *) object, (intptr_t) object->mark(),
134                 object->klass()->external_name());
135         }
136     }
137     return m ;
138 }
139
140 // CASE: neutral 无锁

```

```

141 // TODO-FIXME: for entry we currently inflate and then try to CAS _owner.
142 // If we know we're inflating for entry it's better to inflate by swinging a
143 // pre-locked objectMonitor pointer into the object header. A successful
144 // CAS inflates the object *and* confers ownership to the inflating thread.
145 // In the current implementation we use a 2-step mechanism where we CAS()
146 // to inflate and then CAS() again to try to swing _owner from NULL to Self.
147 // An inflateTry() method that we could call from fast_enter() and slow_enter()
148 // would be useful.
149
150 assert (mark->is_neutral(), "invariant");
151 ObjectMonitor * m = omAlloc (Self) ;
152 // prepare m for installation - set monitor to initial state
153 m->Recycle();
154 m->set_header(mark);
155 m->set_owner(NULL);
156 m->set_object(object);
157 m->OwnerIsThread = 1 ;
158 m->_recursions = 0 ;
159 m->_Responsible = NULL ;
160 m->_SpinDuration = ObjectMonitor::Knob_SpinLimit ; // consider: keep metastats by type/cla
161
162 if (Atomic::cmpxchg_ptr (markOopDesc::encode(m), object->mark_addr(), mark) != mark) {
163     m->set_object (NULL) ;
164     m->set_owner (NULL) ;
165     m->OwnerIsThread = 0 ;
166     m->Recycle() ;
167     omRelease (Self, m, true) ;
168     m = NULL ;
169     continue ;
170     // interference - the markword changed - just retry.
171     // The state-transitions are one-way, so there's no chance of
172     // live-lock -- "Inflated" is an absorbing state.
173 }
174
175 // Hopefully the performance counters are allocated on distinct
176 // cache lines to avoid false sharing on MP systems ...
177 if (ObjectMonitor::_sync_Inflations != NULL) ObjectMonitor::_sync_Inflations->inc() ;
178 TEVENT(Inflate: overwrite neutral) ;
179 if (TraceMonitorInflation) {
180     if (object->is_instance()) {
181         ResourceMark rm;
182         tty->print_cr("Inflating object " INTPTR_FORMAT " , mark " INTPTR_FORMAT " , type %s",
183             (void *) object, (intptr_t) object->mark(),
184             object->klass()->external_name());
185     }
186 }
187 return m ;
188 }
189 }

```



膨胀过程的实现比较复杂, 大概实现过程如下:

- 1、整个膨胀过程在自旋下完成;
- 2、mark->has_monitor() 方法判断当前是否为重量级锁 (上图18-25行), 即Mark Word的锁标识位为 **10**, 如果当前状态为重量级锁, 执行步骤 (4) ;
- 3、mark->monitor() 方法获取指向ObjectMonitor的指针, 并返回, 说明膨胀过程已经完成;
- 4、如果当前锁处于**膨胀中** (上图33-37行), 说明该锁正在被其它线程执行膨胀操作, 则当前线程就进行自旋等待锁膨胀完成, 这里需要注意一点, **它**不会一直占用cpu资源, 每隔一段时间会通过os::NakedYield方法放弃cpu资源, 或通过park方法挂起; 如果其他线程完成锁的膨胀操作, 则退出自旋
- 5、如果当前是**轻量级锁**状态 (上图58-138行), 即锁标识位为 **00**, 膨胀过程如下:

1. 通过omAlloc方法, 获取一个可用的ObjectMonitor monitor, 并重置monitor数据;
2. 通过CAS尝试将Mark Word设置为markOopDesc:INFLATING, 标识当前锁正在膨胀中, 如果CAS失败, 说明同一时刻其它线程已经将Mark Word设置为INFLATING, 当前线程进行自旋等待膨胀完成;
3. 如果CAS成功, 设置monitor的各个字段: _header、_owner和_object等, 并返回;

- 6、如果是**无锁** (中立, 上图150-186行), 重置监视器值;

monitor竞争

当锁膨胀完成并返回对应的monitor时, 并不表示该线程竞争到了锁, 真正的锁竞争发生在ObjectMonitor::enter方法中。



```

1 void ATTR ObjectMonitor::enter(TRAPS) {
2 // The following code is ordered to check the most common cases first
3 // and to reduce RTS->RTO cache line upgrades on SPARC and IA32 processors.
4 Thread * const Self = THREAD ;
5 void * cur ;
6
7 cur = Atomic::cmpxchg_ptr (Self, &_owner, NULL) ;
8 if (cur == NULL) { //CAS成功
9 // Either ASSERT _recursions == 0 or explicitly set _recursions = 0.
10 assert (_recursions == 0 , "invariant") ;
11 assert (_owner == Self, "invariant") ;
12 // CONSIDER: set or assert OwnerIsThread == 1
13 return ;
14 }
15
16 if (cur == Self) { //重入锁

```

```

17 // TODO-FIXME: check for integer overflow! BUGID 6557169.
18 _recursions ++ ;
19 return ;
20 }
21
22 if (Self->is_lock_owned ((address)cur)) {
23     assert (_recursions == 0, "internal state error");
24     _recursions = 1 ;
25     // Commute owner from a thread-specific on-stack BasicLockObject address to
26     // a full-fledged "Thread *".
27     _owner = Self ;
28     OwnerIsThread = 1 ;
29     return ;
30 }
31
32 // We've encountered genuine contention.
33 assert (Self->_Stalled == 0, "invariant") ;
34 Self->_Stalled = intptr_t(this) ;
35
36 // Try one round of spinning *before* enqueueing Self
37 // and before going through the awkward and expensive state
38 // transitions. The following spin is strictly optional ...
39 // Note that if we acquire the monitor from an initial spin
40 // we forgo posting JVMTI events and firing DTRACE probes.
41 if (Knob_SpinEarly && TrySpin (Self) > 0) {
42     assert (_owner == Self , "invariant") ;
43     assert (_recursions == 0 , "invariant") ;
44     assert (((oop)(object()))->mark() == markOopDesc::encode(this), "invariant") ;
45     Self->_Stalled = 0 ;
46     return ;
47 }
48
49 assert (_owner != Self , "invariant") ;
50 assert (_succ != Self , "invariant") ;
51 assert (Self->is_Java_thread() , "invariant") ;
52 JavaThread * jt = (JavaThread *) Self ;
53 assert (!SafepointSynchronize::is_at_safepoint(), "invariant") ;
54 assert (jt->thread_state() != _thread_blocked , "invariant") ;
55 assert (this->object() != NULL , "invariant") ;
56 assert (_count >= 0, "invariant") ;
57
58 // Prevent deflation at STW-time. See deflate_idle_monitors() and is_busy().
59 // Ensure the object-monitor relationship remains stable while there's contention.
60 Atomic::inc_ptr(&_count);
61
62 EventJavaMonitorEnter event;
63
64 { // Change java thread status to indicate blocked on monitor enter.
65     JavaThreadBlockedOnMonitorEnterState jthmes(jt, this);
66
67     DTRACE_MONITOR_PROBE(contended__enter, this, object(), jt);
68     if (JvmtiExport::should_post_monitor_contended_enter()) {
69         JvmtiExport::post_monitor_contended_enter(jt, this);
70     }
71
72     OSThreadContendState osts(Self->osthread());
73     ThreadBlockInVM tbvm(jt);
74
75     Self->set_current_pending_monitor(this);
76
77     // TODO-FIXME: change the following for(;;) loop to straight-line code.
78     for (;;) {
79         jt->set_suspend_equivalent();
80         // cleared by handle_special_suspend_equivalent_condition()
81         // or java_suspend_self()
82
83         EnterI (THREAD) ;
84
85         ...省略...139 }

```



- 1、通过CAS尝试把monitor的_owner字段设置为当前线程；
- 2、如果设置之前的_owner指向当前线程，说明当前线程再次进入monitor，即重入锁，执行_recursions ++，记录重入的次数；
- 3、如果之前的_owner指向的地址在当前线程中，这种描述有点拗口，换一种说法：之前_owner指向的BasicLock在当前线程栈上，说明当前线程是monitor，设置_recursions为1，_owner为当前线程，该线程成功获得锁并返回；
- 4、如果获取锁失败，则等待锁的释放；

monitor等待

monitor竞争失败的线程，通过自旋执行ObjectMonitor::EnterI方法等待锁的释放，EnterI方法的部分逻辑实现如下：



```

1 ObjectWaiter node(Self) ;
2 Self->_ParkEvent->reset() ;
3 node._prev = (ObjectWaiter *) 0xBAD ;
4 node.TState = ObjectWaiter::TS_CXQ ;
5

```

```

6    // Push "Self" onto the front of the _cxq.
7    // Once on cxq/EntryList, Self stays on-queue until it acquires the lock.
8    // Note that spinning tends to reduce the rate at which threads
9    // enqueue and dequeue on EntryList|cxq.
10   ObjectWaiter * nxt ;
11   for (;;) {
12       node._next = nxt = _cxq ;
13       if (Atomic::cmpxchg_ptr (&node, &_cxq, nxt) == nxt) break ;
14
15       // Interference - the CAS failed because _cxq changed. Just retry.
16       // As an optional optimization we retry the lock.
17       if (TryLock (Self) > 0) {
18           assert (_succ != Self , "invariant") ;
19           assert (_owner == Self , "invariant") ;
20           assert (_Responsible != Self , "invariant") ;
21           return ;
22       }
23   }

```



- 1、当前线程被封装成ObjectWaiter对象node，状态设置成ObjectWaiter::TS_CXQ；
- 2、在for循环中，通过CAS把node节点push到_cxq列表中，同一时刻可能有多个线程把自己的node节点push到_cxq列表中；
- 3、node节点push到_cxq列表之后，通过自旋尝试获取锁，如果还是没有获取到锁，则通过park将当前线程挂起，等待被唤醒，实现如下：



```

1  for (;;) {
2
3      if (TryLock (Self) > 0) break ;
4      assert (_owner != Self, "invariant") ;
5
6      if ((SyncFlags & 2) && _Responsible == NULL) {
7          Atomic::cmpxchg_ptr (Self, &_Responsible, NULL) ;
8      }
9
10     // park self
11     if (_Responsible == Self || (SyncFlags & 1)) {
12         TEVENT (Inflated enter - park TIMED) ;
13         Self->_ParkEvent->park ((jlong) RecheckInterval) ;
14         // Increase the RecheckInterval, but clamp the value.
15         RecheckInterval *= 8 ;
16         if (RecheckInterval > 1000) RecheckInterval = 1000 ;
17     } else {
18         TEVENT (Inflated enter - park UNTIMED) ;
19         Self->_ParkEvent->park() ; //当前线程挂起
20     }
21
22     if (TryLock(Self) > 0) break ;
23
24     // The lock is still contested.
25     // Keep a tally of the # of futile wakeups.
26     // Note that the counter is not protected by a lock or updated by atomics.
27     // That is by design - we trade "lossy" counters which are exposed to
28     // races during updates for a lower probe effect.
29     TEVENT (Inflated enter - Futile wakeup) ;
30     if (ObjectMonitor::_sync_FutileWakeups != NULL) {
31         ObjectMonitor::_sync_FutileWakeups->inc() ;
32     }
33     ++ nWakeups ;
34
35     // Assuming this is not a spurious wakeup we'll normally find _succ == Self.
36     // We can defer clearing _succ until after the spin completes
37     // TrySpin() must tolerate being called with _succ == Self.
38     // Try yet another round of adaptive spinning.
39     if ((Knob_SpinAfterFutile & 1) && TrySpin (Self) > 0) break ;
40
41     // We can find that we were unpark()ed and redesignated _succ while
42     // we were spinning. That's harmless. If we iterate and call park(),
43     // park() will consume the event and return immediately and we'll
44     // just spin again. This pattern can repeat, leaving _succ to simply
45     // spin on a CPU. Enable Knob_ResetEvent to clear pending unparks().
46     // Alternately, we can sample fired() here, and if set, forgo spinning
47     // in the next iteration.
48
49     if ((Knob_ResetEvent & 1) && Self->_ParkEvent->fired()) {
50         Self->_ParkEvent->reset() ;
51         OrderAccess::fence() ;
52     }
53     if (_succ == Self) _succ = NULL ;
54
55     // Invariant: after clearing _succ a thread *must* retry _owner before parking.
56     OrderAccess::fence() ;
57 }

```



- 4、当该线程被唤醒时，会从挂起的点继续执行，通过ObjectMonitor::TryLock尝试获取锁，TryLock方法实现如下：



```

1 int ObjectMonitor::TryLock (Thread * Self) {
2     for (;;) {
3         void * own = _owner ;
4         if (own != NULL) return 0 ;
5         if (Atomic::cmpxchg_ptr (Self, &_owner, NULL) == NULL) { //CAS成功, 获取锁
6             // Either guarantee _recursions == 0 or set _recursions = 0.
7             assert (_recursions == 0, "invariant") ;
8             assert (_owner == Self, "invariant") ;
9             // CONSIDER: set or assert that OwnerIsThread == 1
10            return 1 ;
11        }
12        // The lock had been free momentarily, but we lost the race to the lock.
13        // Interference -- the CAS failed.
14        // We can either return -1 or retry.
15        // Retry doesn't make as much sense because the lock was just acquired.
16        if (true) return -1 ;
17    }
18 }

```



其本质就是通过CAS设置monitor的_owner字段为当前线程, 如果CAS成功, 则表示该线程获取了锁, 跳出自旋操作, 执行同步代码, 否则继续被挂起。

monitor释放

当某个持有锁的线程执行完同步代码块时, 会进行锁的释放, 给其它线程机会执行同步代码, 在HotSpot中, 通过退出monitor的方式实现锁的释放, 具体实现位于ObjectMonitor::exit方法中。



```

1 void ATTR ObjectMonitor::exit(bool not_suspended, TRAPS) {
2     Thread * Self = THREAD ;
3     if (THREAD != _owner) {
4         if (THREAD->is_lock_owned((address) _owner)) {
5             // Transmute _owner from a BasicLock pointer to a Thread address.
6             // We don't need to hold _mutex for this transition.
7             // Non-null to Non-null is safe as long as all readers can
8             // tolerate either flavor.
9             assert (_recursions == 0, "invariant") ;
10            _owner = THREAD ;
11            _recursions = 0 ;
12            OwnerIsThread = 1 ;
13        } else {
14            // NOTE: we need to handle unbalanced monitor enter/exit
15            // in native code by throwing an exception.
16            // TODO: Throw an IllegalMonitorStateException ?
17            TEVENT (Exit - Throw IMSX) ;
18            assert(false, "Non-balanced monitor enter/exit!");
19            if (false) {
20                THROW(vmSymbols::java_lang_IllegalMonitorStateException());
21            }
22            return;
23        }
24    }
25
26    if (_recursions != 0) {
27        _recursions--; // this is simple recursive enter
28        TEVENT (Inflated exit - recursive) ;
29        return ;
30    }
31    ...省略...

```



1、如果是重量级锁的释放, monitor中的_owner指向当前线程, 即THREAD == _owner;

2、根据不同的策略(由QMode指定), 从cxq或EntryList中获取头节点, 通过ObjectMonitor::ExitEpilog方法唤醒该节点封装的线程, 唤醒操作, 实现如下:



```

1 void ObjectMonitor::ExitEpilog (Thread * Self, ObjectWaiter * Wakee) {
2     assert (_owner == Self, "invariant") ;
3
4     // Exit protocol:
5     // 1. ST_succ = wakee
6     // 2. membar #loadstore|#storestore;
7     // 2. ST_owner = NULL
8     // 3. unpark(wakee)
9
10    _succ = Knob_SuccEnabled ? Wakee->_thread : NULL ;
11    ParkEvent * Trigger = Wakee->_event ;
12
13    // Hygiene -- once we've set _owner = NULL we can't safely dereference Wakee again.
14    // The thread associated with Wakee may have grabbed the lock and "Wakee" may be
15    // out-of-scope (non-extant).
16    Wakee = NULL ;
17

```



```
18 // Drop the lock
19 OrderAccess::release_store_ptr (&_owner, NULL) ;
20 OrderAccess::fence() ; // ST _owner vs LD in unpark()
21
22 if (SafepointSynchronize::do_call_back()) {
23     TEVENT (unpark before SAFEPOINT) ;
24 }
25
26 DTRACE_MONITOR_PROBE(contended__exit, this, object(), Self);
27 Trigger->unpark() ;
28
29 // Maintain stats and report events to JVMTI
30 if (ObjectMonitor::_sync_Parks != NULL) {
31     ObjectMonitor::_sync_Parks->inc() ;
32 }
33 }
```

3、被唤醒的线程，继续执行monitor的竞争；

四.总结

本文重点介绍了Synchronized原理以及JVM对Synchronized的优化。简单来说解决三种场景：

- 1) 只有一个线程进入临界区，偏向锁
- 2) 多个线程交替进入临界区，轻量级锁
- 3) 多线程同时进入临界区，重量级锁

参考：

《深入理解 Java 虚拟机》

JVM源码分析之synchronized实现

----- 掌控自己的生命轨迹，身心自由！ -----

标签: jdk源码剖析

好文要顶

关注我

收藏该文

只会一点java

关注 - 6

粉丝 - 46

+加关注

« 上一篇: jdk源码剖析一：OpenJDK-Hotspot源码包目录结构
» 下一篇: jdk源码剖析四：JDK1.7升级1.8 HashMap原理的变化

Feedback

#1楼 2017-11-30 00:19 by christoph
老哥牛逼啊，这么牛逼的文章居然没有看啊，
#2楼[楼主] 2017-11-30 12:16 by 只会一点java
甚是欣慰
#3楼 2017-12-27 19:00 by Mr King~
感谢您的分享。对图1有些不太明白的地方。图一中，线程2，进行CAS MarkWord替换，什么时候会出现不成功的情况，什么时候会出现成功的情况？